

CAP5415-Computer Vision

Lecture 3-Edge Maps and Histograms

Ulas Bagci

bagci@ucf.edu

Outline

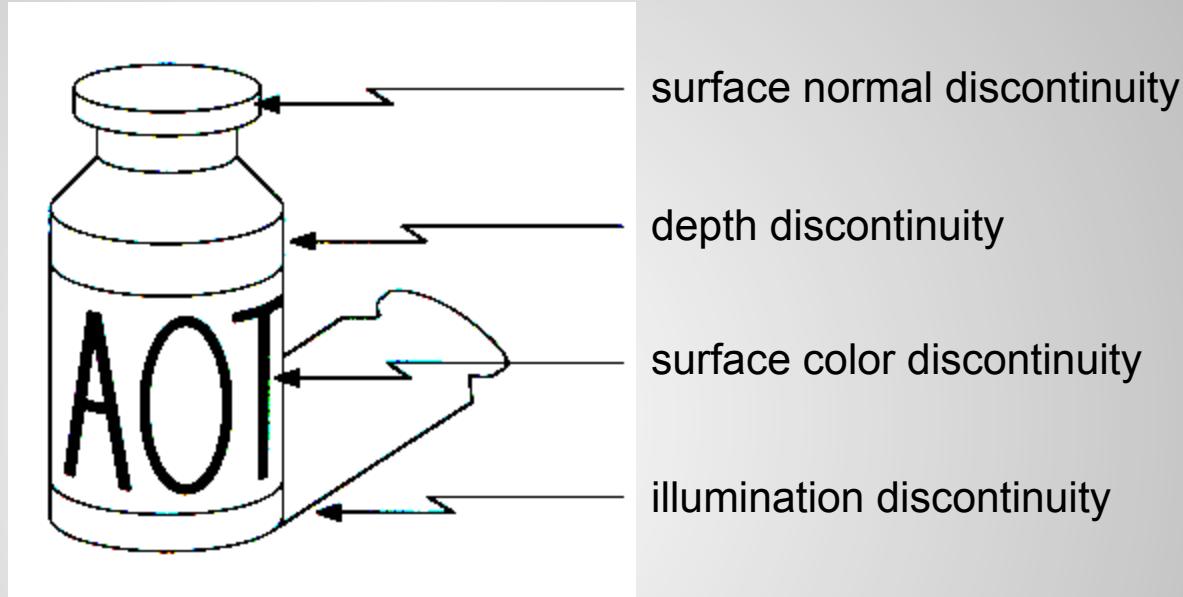
- Continue edge detection and filtering methods,
 - Histogram-based analysis
-
- *Read Szeliski, Chapter 3.*
 - *Read Shah, Chapter 2.*
 - *Read/Program CV with Python, Chapters 1 and 2.*

Revisiting Edge Detection

- **Goal:** Identify sudden changes (discontinuities) in an image
 - Most semantic and shape information from the image can be encoded in the edges
 - More compact than pixels
 - Marks the border of an object

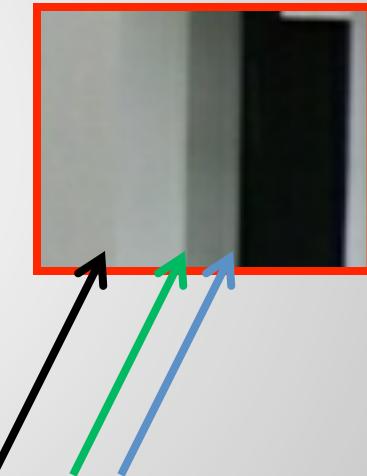


Origins of Edges

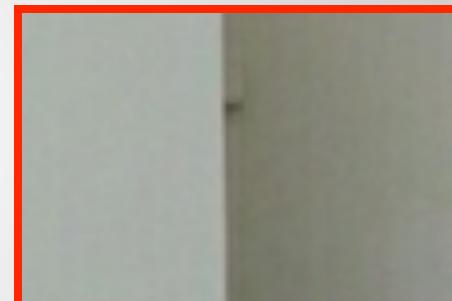


- Edges are caused by a variety of factors

Close-up Edges

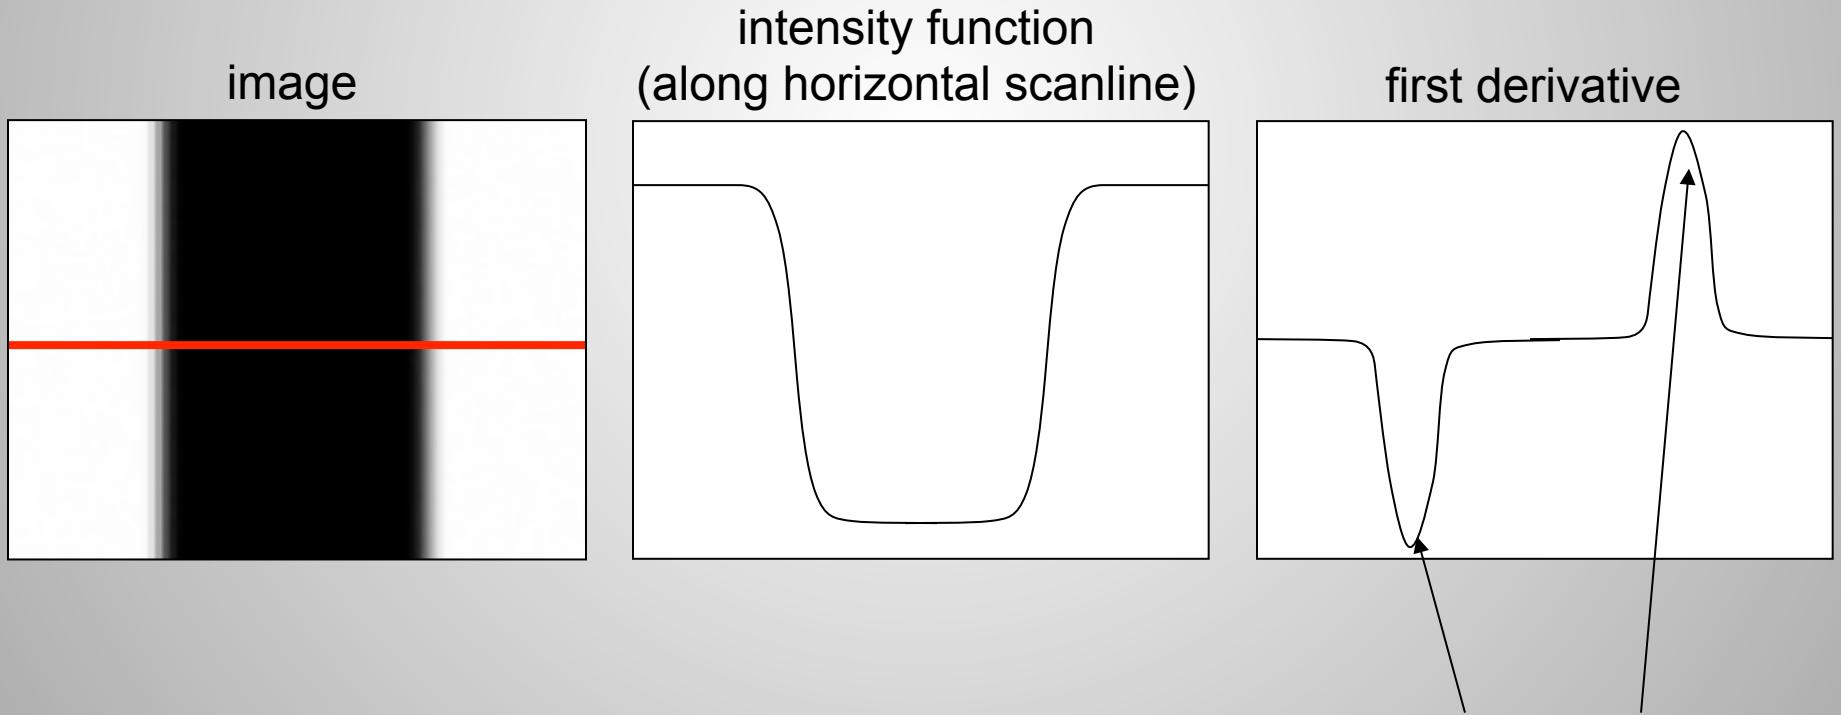


Close-up Edges



Characterizing Edges

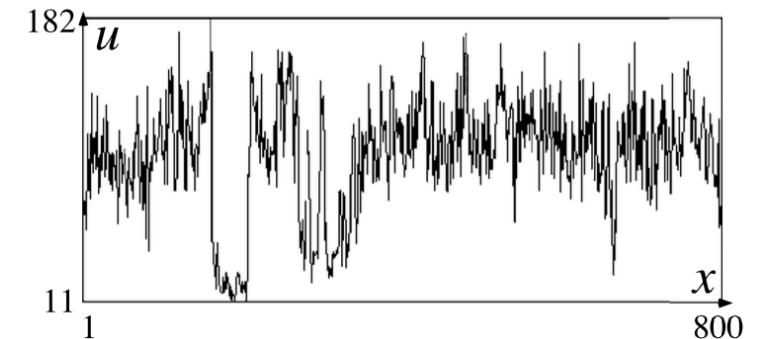
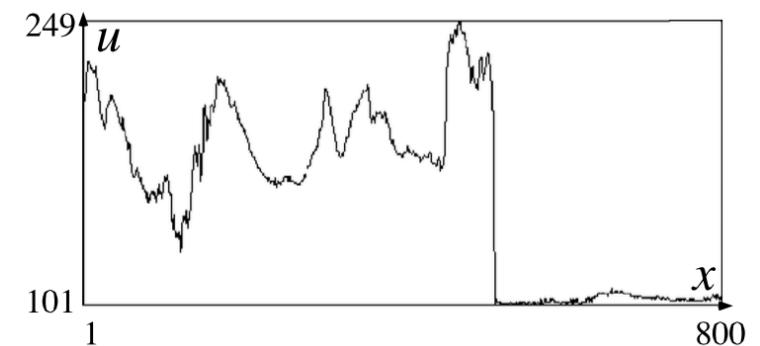
- An edge is a place of rapid change in the image intensity function



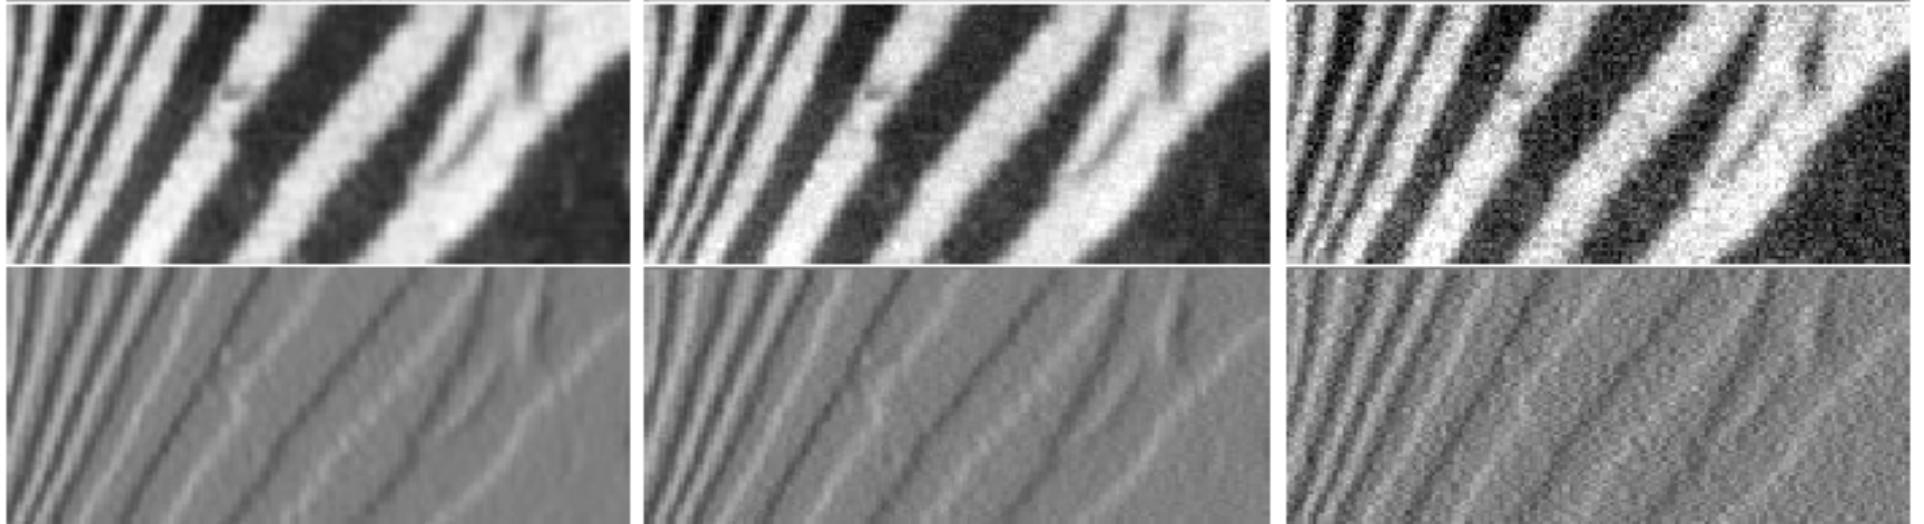
Slide Credit: James Hays

edges correspond to
extrema of derivative

Intensity Profile



Effects of Noise



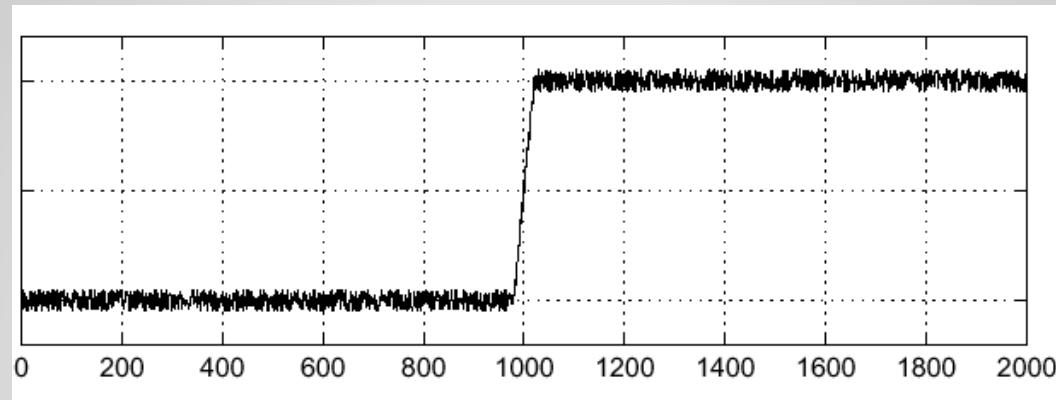
Increasing noise



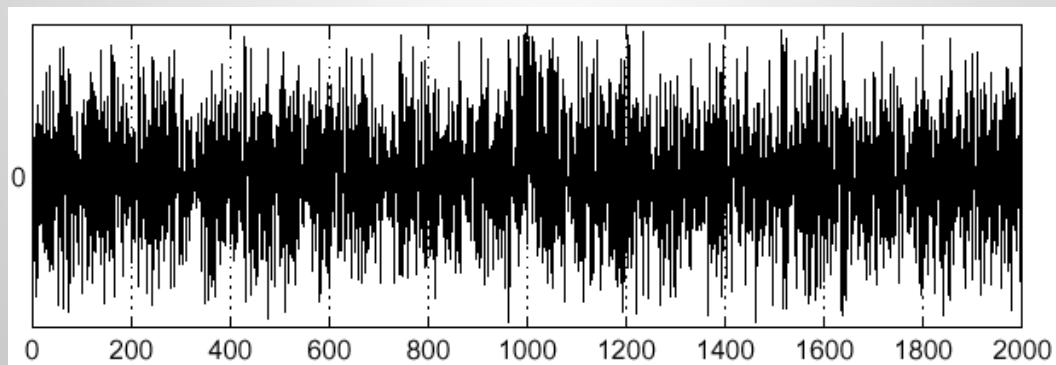
Zero mean additive Gaussian noise

Effects of Noise

$f(x)$

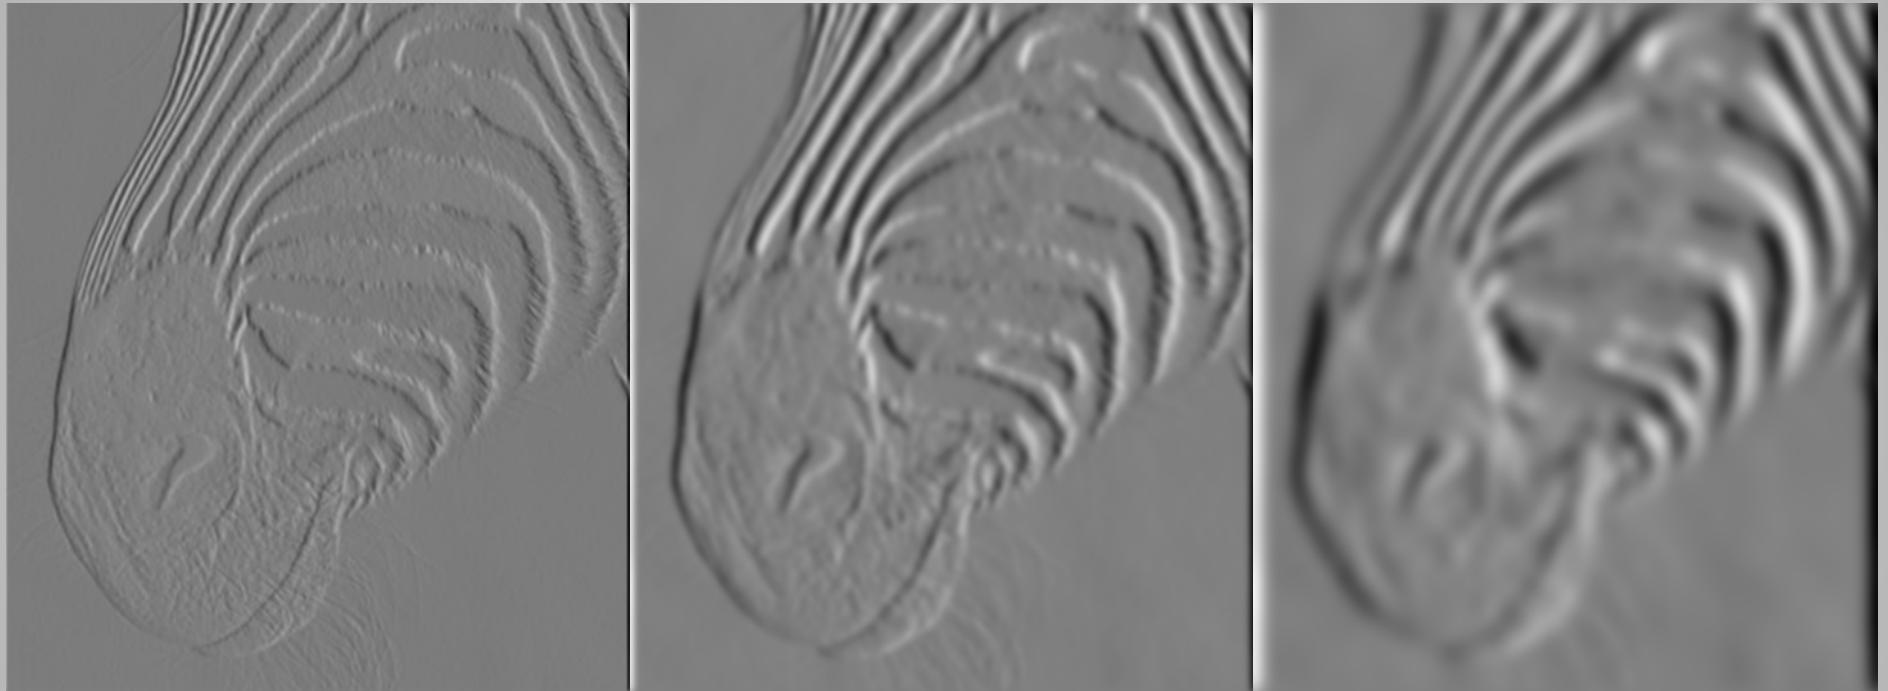


$\frac{d}{dx}f(x)$



Where is the edge?

Solution: Smoothing



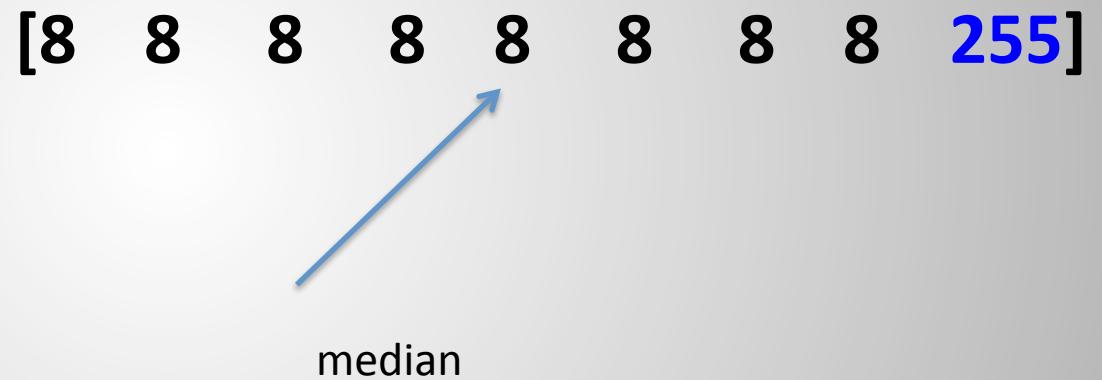
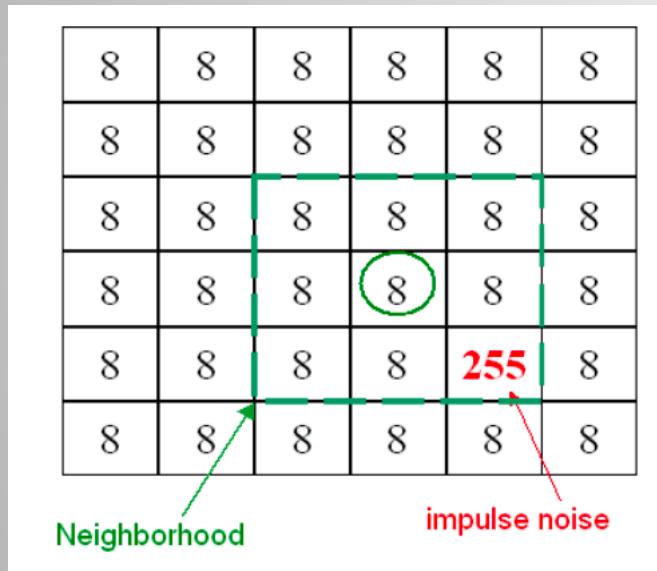
1 pixel

3 pixels

7 pixels

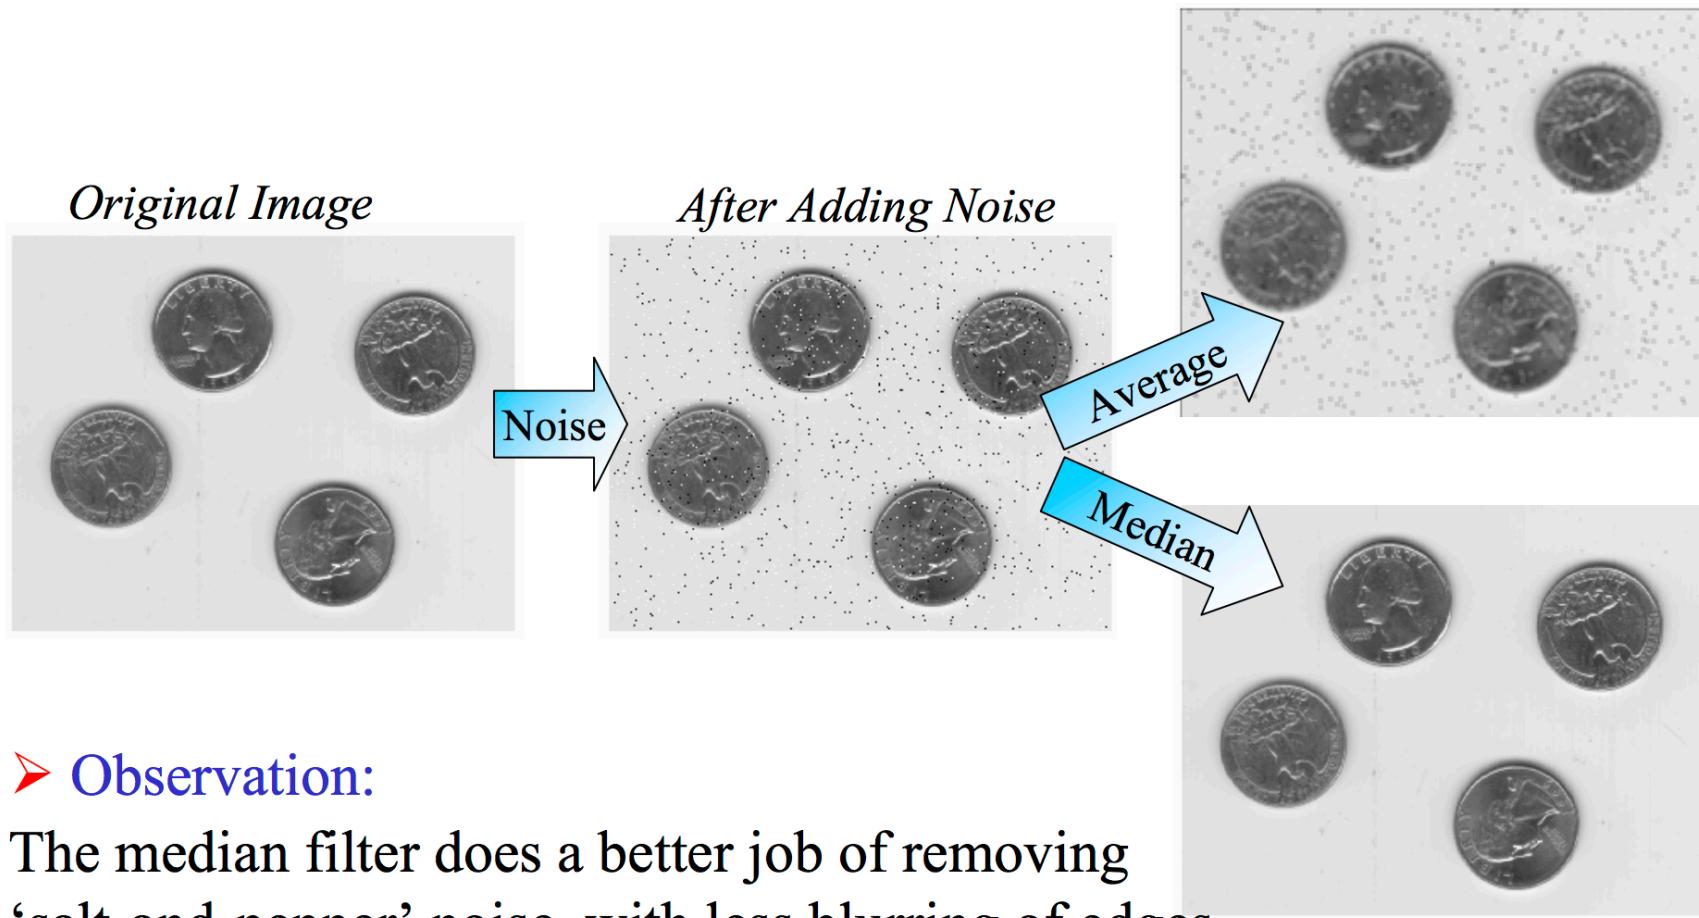
- Smoothing removes noise, but **blurs** edge.

Revisiting Median Filtering



Revisiting Median Filtering

- Original image “*Eight.tif*” with added ‘salt-and-pepper’ noise then filtered with a (3-by-3) averaging filter and a (3-by-3) median filter.



Revisit: Image Gradient

- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

The first diagram shows a horizontal gradient with a red arrow pointing right, labeled $\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$. The second diagram shows a vertical gradient with a red arrow pointing down, labeled $\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$. The third diagram shows a diagonal gradient with a red arrow pointing up and to the right at an angle θ , labeled $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$.

The gradient points in the direction of most rapid increase in intensity

The gradient direction is given by $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$

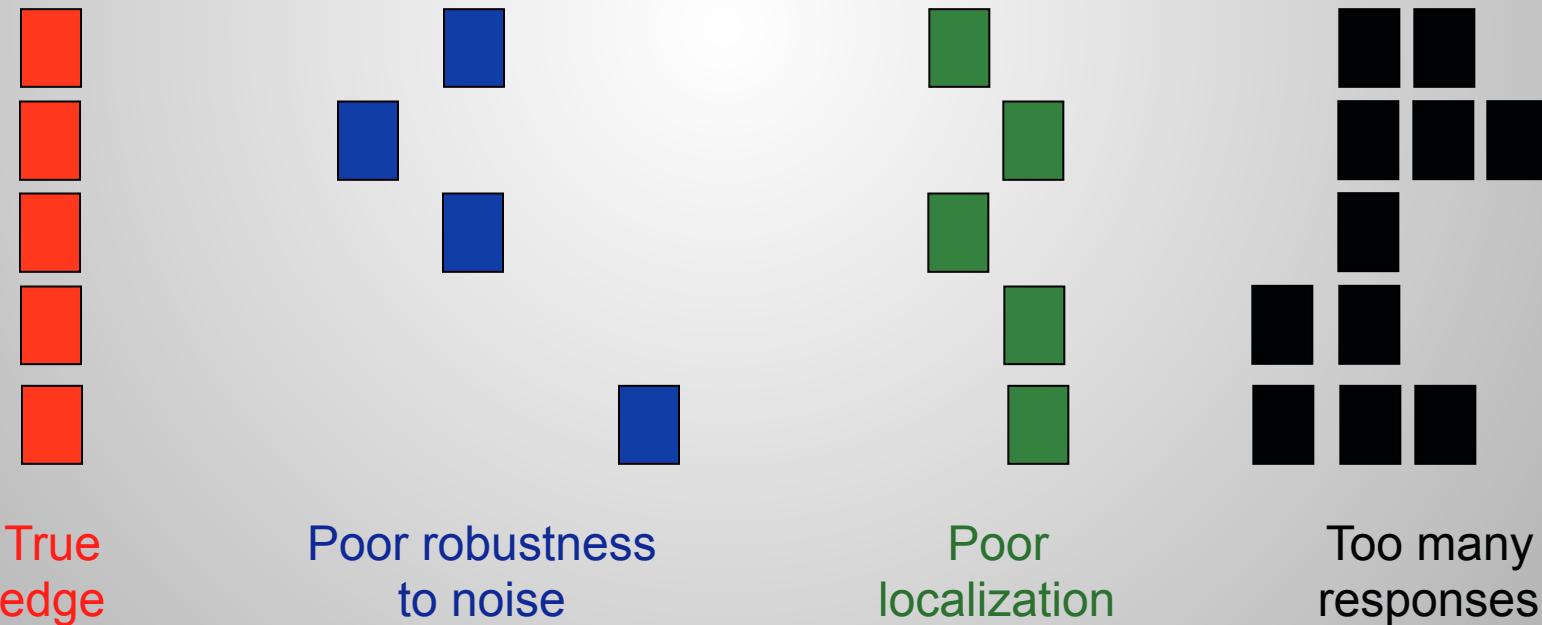
- how does this relate to the direction of the edge?

The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Design Criteria for Edge Detection Problems

- **Good Detection:** minimize prob. of FP (detecting spurious edges) and FN (missing real edges)
- **Good Localization:** must be as close as possible to the true edges
- **Single Response:** must return one point only for each true edge point

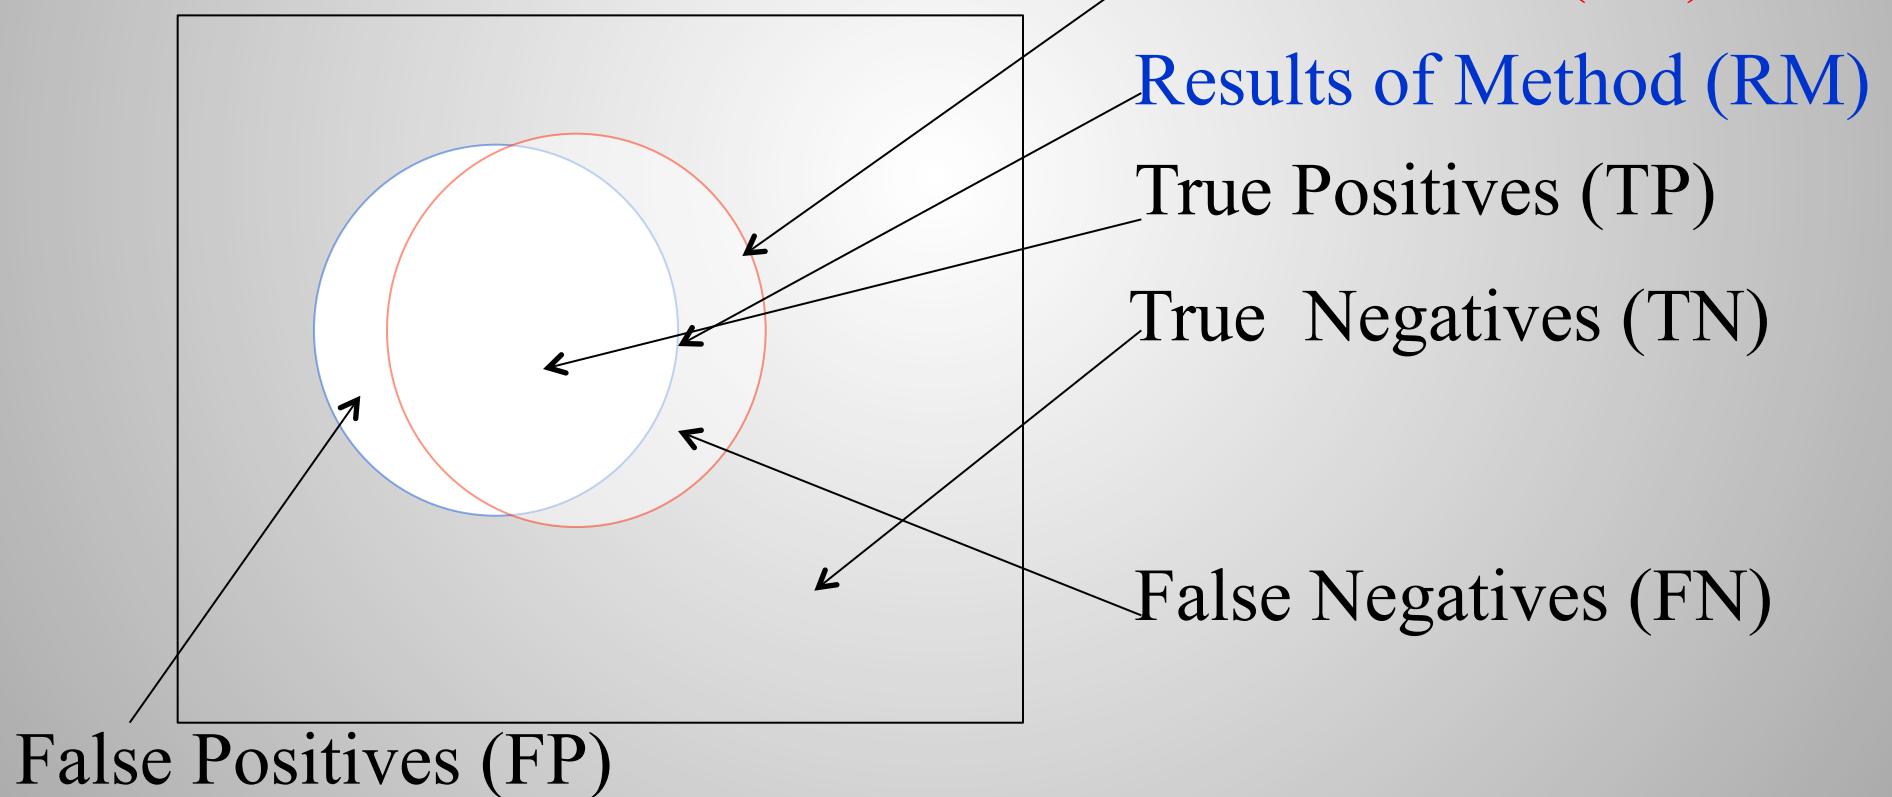


Evaluate Edge Detection

$$\text{precision} = \frac{\text{GT} \cap \text{RM}}{\text{RM}} = \frac{\text{TP}}{\text{RM}}$$

$$\text{recall} = \frac{\text{GT} \cap \text{RM}}{\text{GT}} = \frac{\text{TP}}{\text{GT}}$$

$$F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$



Basic Comparisons of Edge Operators

Gradient:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

Good Localization
Noise Sensitive
Poor Detection

Roberts (2 x 2):

0	1
-1	0

1	0
0	-1

Sobel (3 x 3):

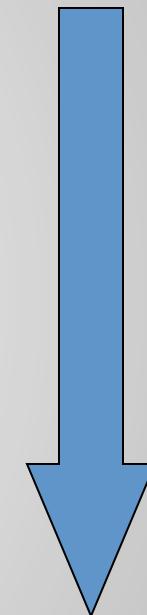
-1	0	1
-1	0	1
-1	0	1

1	1	1
0	0	0
-1	-1	1

Sobel (5 x 5):

-1	-2	0	2	1
-2	-3	0	3	2
-3	-5	0	5	3
-2	-3	0	3	2
-1	-2	0	2	1

1	2	3	2	1
2	3	5	3	2
0	0	0	0	0
-2	-3	-5	-3	-2
-1	-2	-3	-2	-1



Poor Localization
Less Noise Sensitive
Good Detection

Example: Laplacian of Gaussian (LoG) and Canny Edge Detector

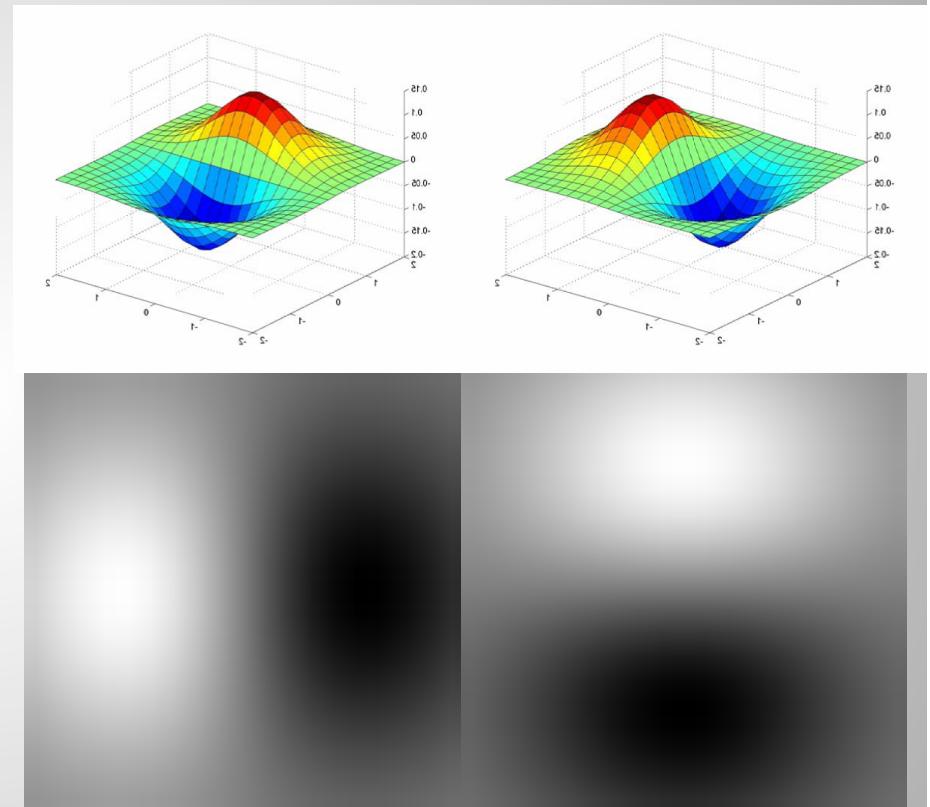
Marr and Hildreth Filtering, 1980.

- Smooth Image with Gaussian Filter
- Applying the Laplacian for a Gaussian-filtered image can be done in one step of convolution.
- Find zero-crossings
- Find slope of zero-crossings
- Apply threshold to slope and mark edges

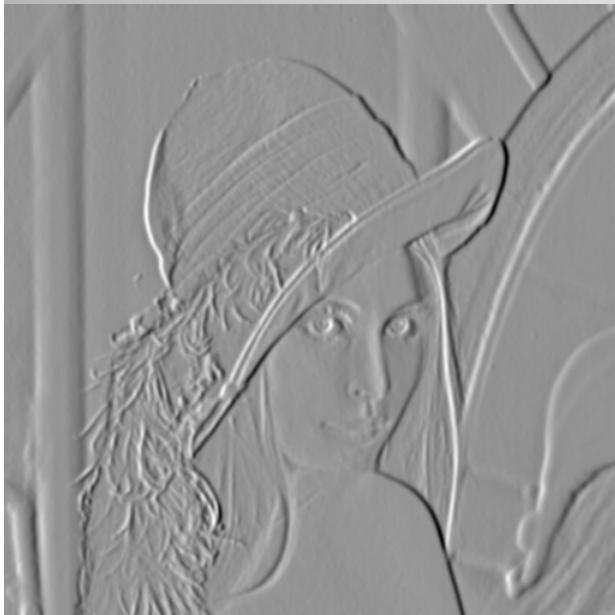
J. Canny. 1986

- Smooth Image with Gaussian filter
- Compute Derivative of filtered image
- Find Magnitude and Orientation of gradient
- Apply Non-max suppression
- Apply Thresholding (Hysteresis)

Example: Canny



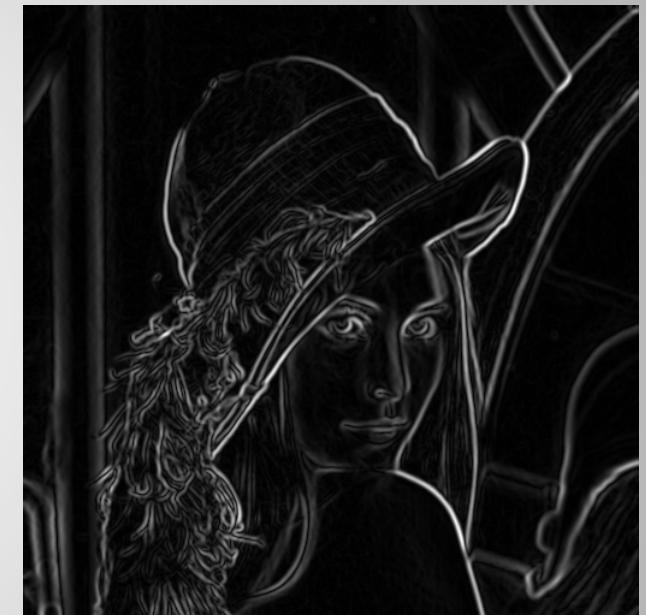
Example: Canny-Gradients



X-Derivative of Gaussian



Y-Derivative of Gaussian



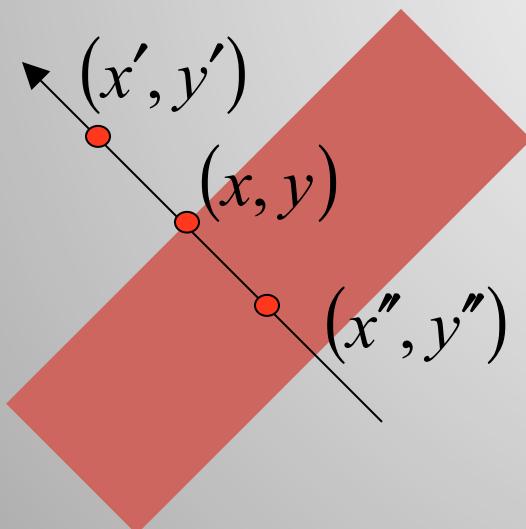
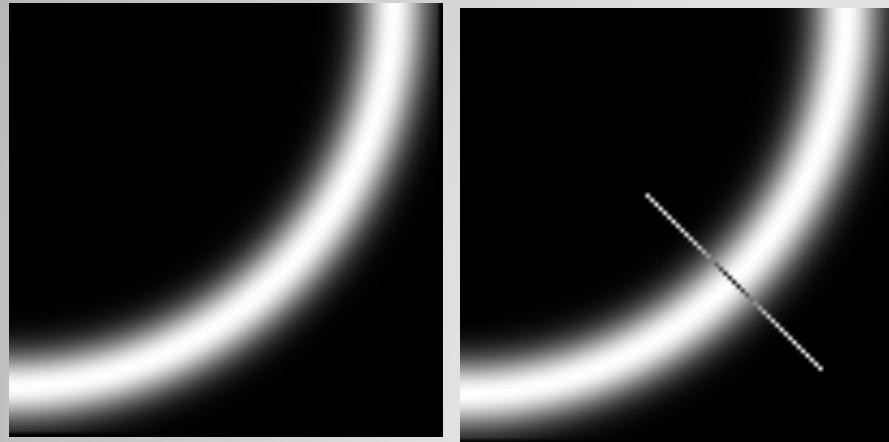
Gradient Magnitude

Example: Gradient Orientation



$$\theta = \text{atan2}(g_y, g_x)$$

Example: Non-maximum suppression



$$M(x, y) = \begin{cases} |\nabla S|(x, y) & \text{if } |\nabla S|(x, y) > |\Delta S|(x', y') \\ & \& |\Delta S|(x, y) > |\Delta S|(x'', y'') \\ 0 & \text{otherwise} \end{cases}$$

x' and x'' are the neighbors of x along normal direction to an edge

Example: Non-maximum suppression



Before Non-Max Suppression



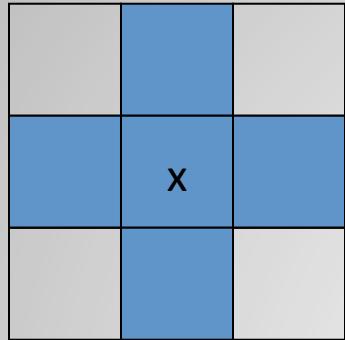
After Non-Max Suppression

Example: Hysteresis Thresholding [L, H]

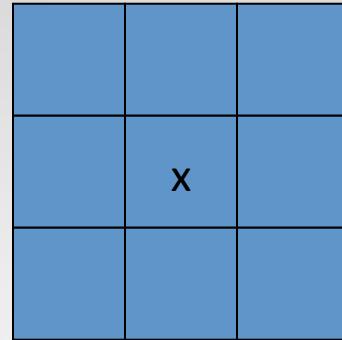
- If the gradient at a pixel is
 - above “**High**”, declare it as an ‘**edge pixel**’
 - below “**Low**”, declare it as a “**non-edge-pixel**”
 - **between** “low” and “high”
 - Consider its neighbors iteratively then declare it an “edge pixel” if it is **connected** to an ‘edge pixel’ **directly** or via pixels **between** “low” and “high”.



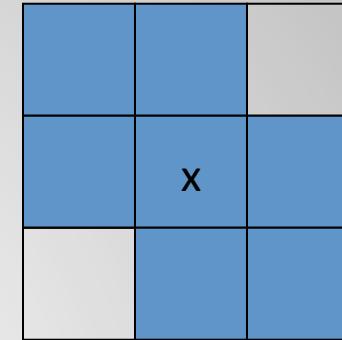
Example: Hysteresis Thresholding [L, H]



4 connected



8 connected



6 connected

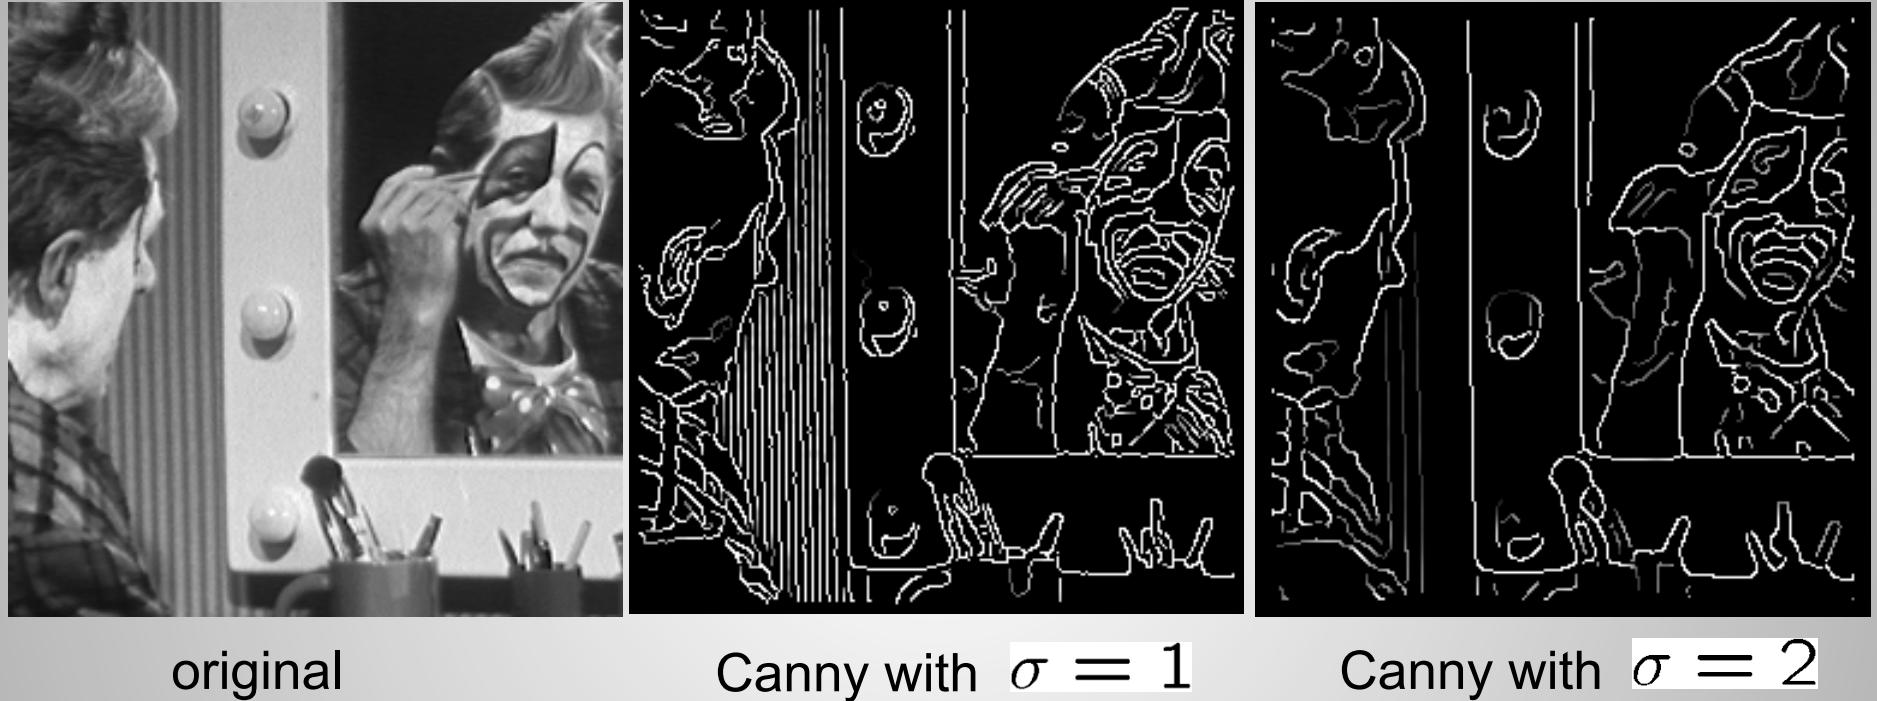
1. Threshold at low/high levels to get weak/strong edge pixels
2. Do connected components, starting from strong edge pixels



Example: Final Canny Edges



Effect of Gaussian Kernel (smoothing)



original

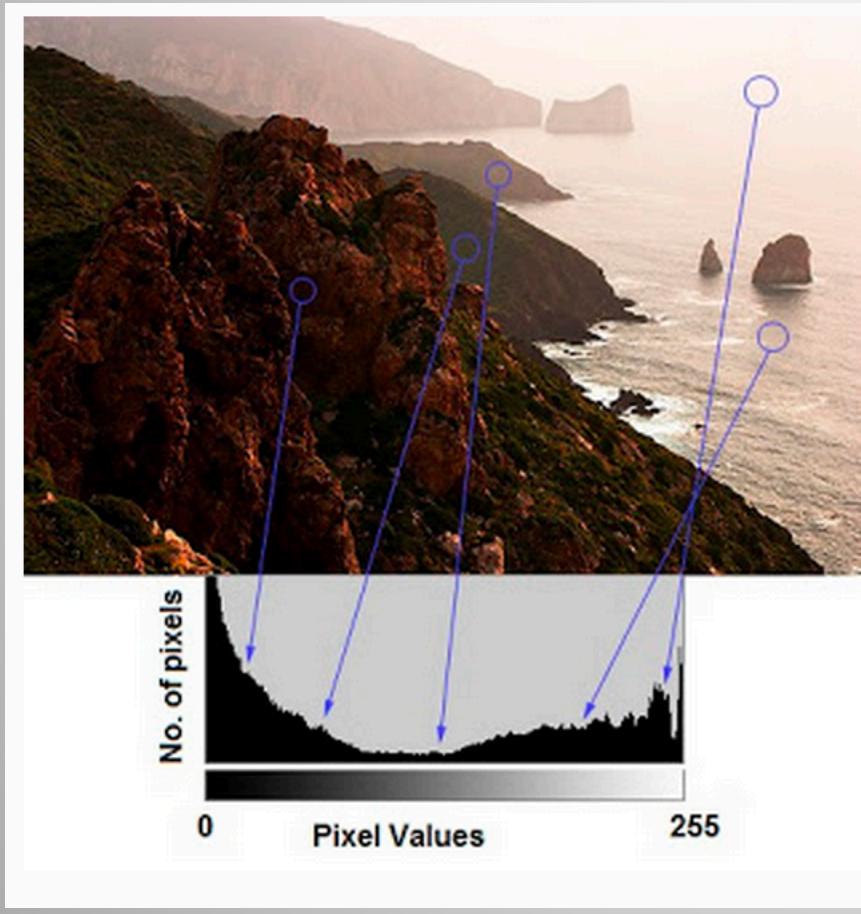
Canny with $\sigma = 1$ Canny with $\sigma = 2$

The choice of depends on desired behavior

- large detects large scale edges
- small detects fine features

Image Histogram

- A **histogram** represents tabulated frequencies, typically by using bars in a graphical diagram.



- It provides a natural bridge between Images and a probabilistic description.

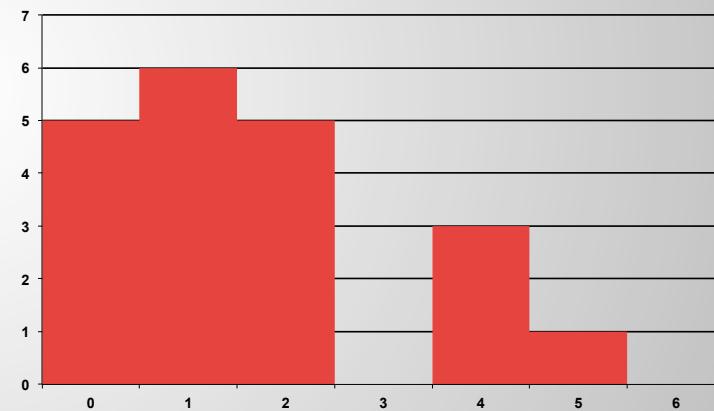
Ex: What is the probability of pixel p at location (x,y) has brightness z ?

- Histogram of a digital image typically has many local minima and maxima, which may complicate its further Processing.
- This problem can be solved by local smoothing of the histogram.

Histogram

0	1	1	2	4
2	1	0	0	2
5	2	0	0	4
1	1	2	4	1

image



histogram

Histogram and PDF

- Assume a scalar image, A, and its histogram H.

$$h(u) = H(u)/|\Omega|$$

- Denominator is the size of histogram (num. of pixels)
- $h \rightarrow$ PDF, relative frequencies are set between 0 and 1.

Entropy / Histogram

- If a probability density h is known, then image information content can be estimated regardless of its interpretation using entropy E .

$$E(X) = - \sum_{k=1}^n h(x_k) \log h(x_k)$$

- Shannon, 1948 [Information entropy].
- Amount of uncertainty about an event associated with a given PDF

Gradation Functions

- When recording image data, there are often particular problems: **lighting, motion blur, noise,...**
 - Uniform illumination is desired
 - Smoothing (denoising)
 - Sharpening

Gradation Functions

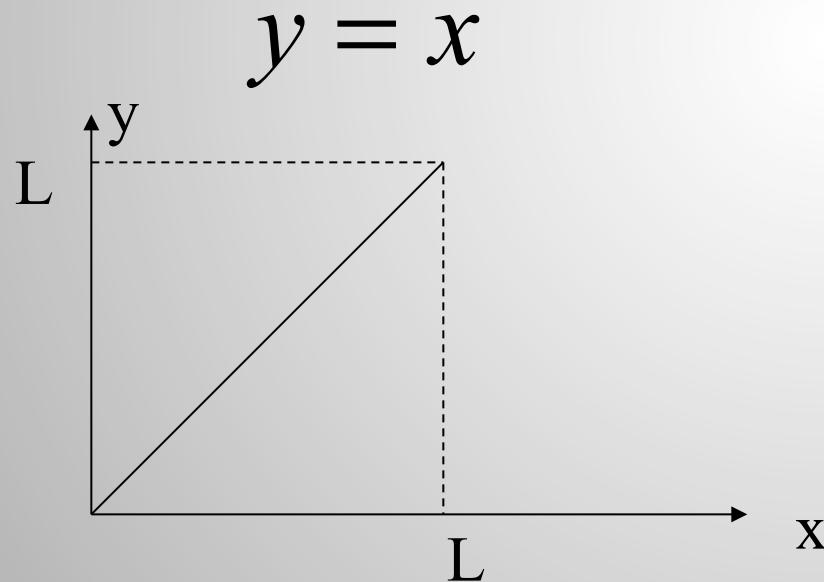
- When recording image data, there are often particular problems: **lighting, motion blur, noise,...**
 - Uniform illumination is desired
 - Smoothing (denoising)
 - Sharpening

Gradation Functions

- We transform an image A into a new image A_{new} of the same size, by mapping a grey level u at pixel location p in A by a gradation function g onto a grey level $v=g(u)$.

Gradation Functions

- We transform an image A into a new image A_{new} of the same size, by mapping a grey level u at pixel location p in A by a gradation function g onto a grey level $v=g(u)$.

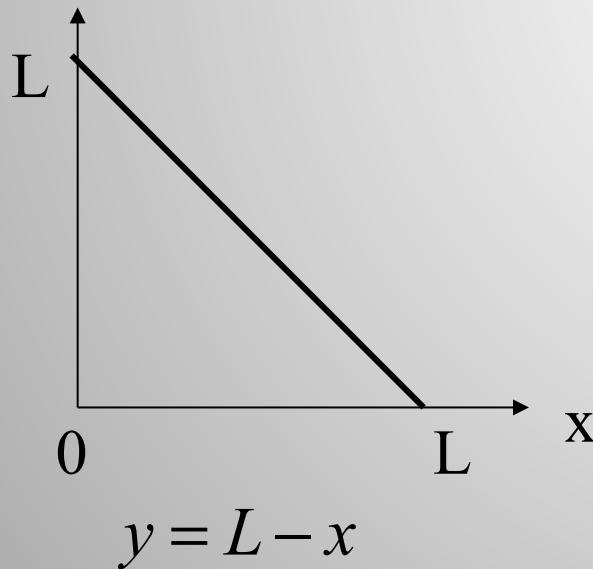


No influence on visual quality at all

Gradation Functions

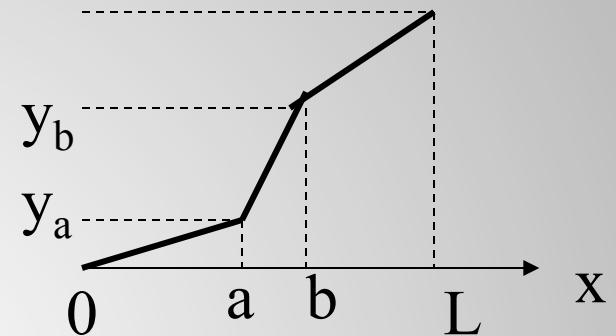
- We transform an image A into a new image A_{new} of the same size, by mapping a grey level u at pixel location p in A by a gradation function g onto a grey level $v=g(u)$.

- Digital negative



Gradation Functions-**Contrast Stretching**

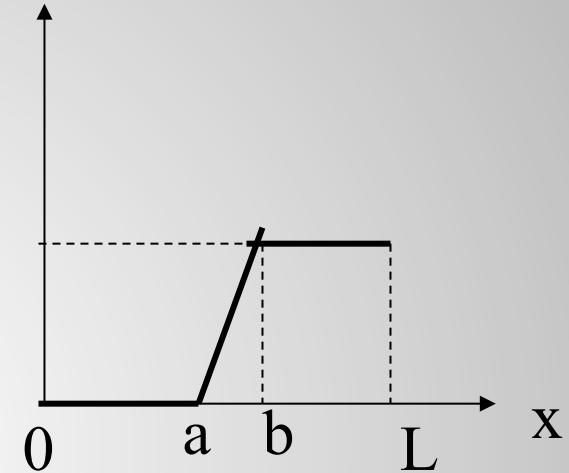
$$y = \begin{cases} \alpha x & 0 \leq x < a \\ \beta(x-a) + y_a & a \leq x < b \\ \gamma(x-b) + y_b & b \leq x < L \end{cases}$$



$$a = 50, b = 150, \alpha = 0.2, \beta = 2, \gamma = 1, y_a = 30, y_b = 200$$

Gradation Functions-Clipping

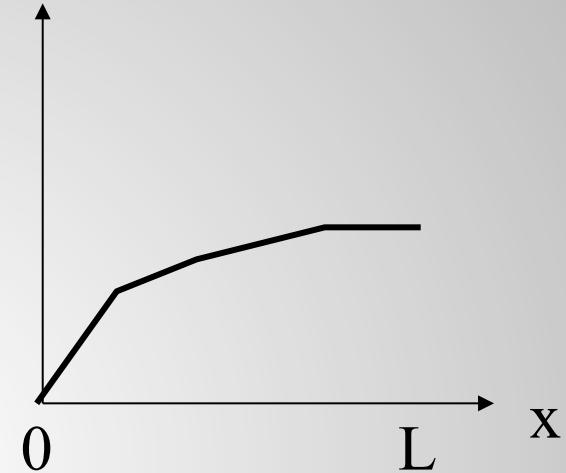
$$y = \begin{cases} 0 & 0 \leq x < a \\ \beta(x-a) & a \leq x < b \\ \beta(b-a) & b \leq x < L \end{cases}$$



$$a = 50, b = 150, \beta = 2$$

Gradation Functions-Range Compression

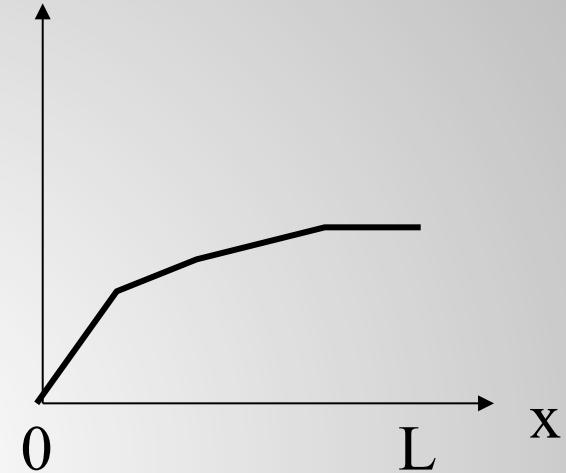
$$y = c \log_{10}(1 + x)$$



$$c=100$$

Gradation Functions-Range Compression

$$y = c \log_{10}(1 + x)$$

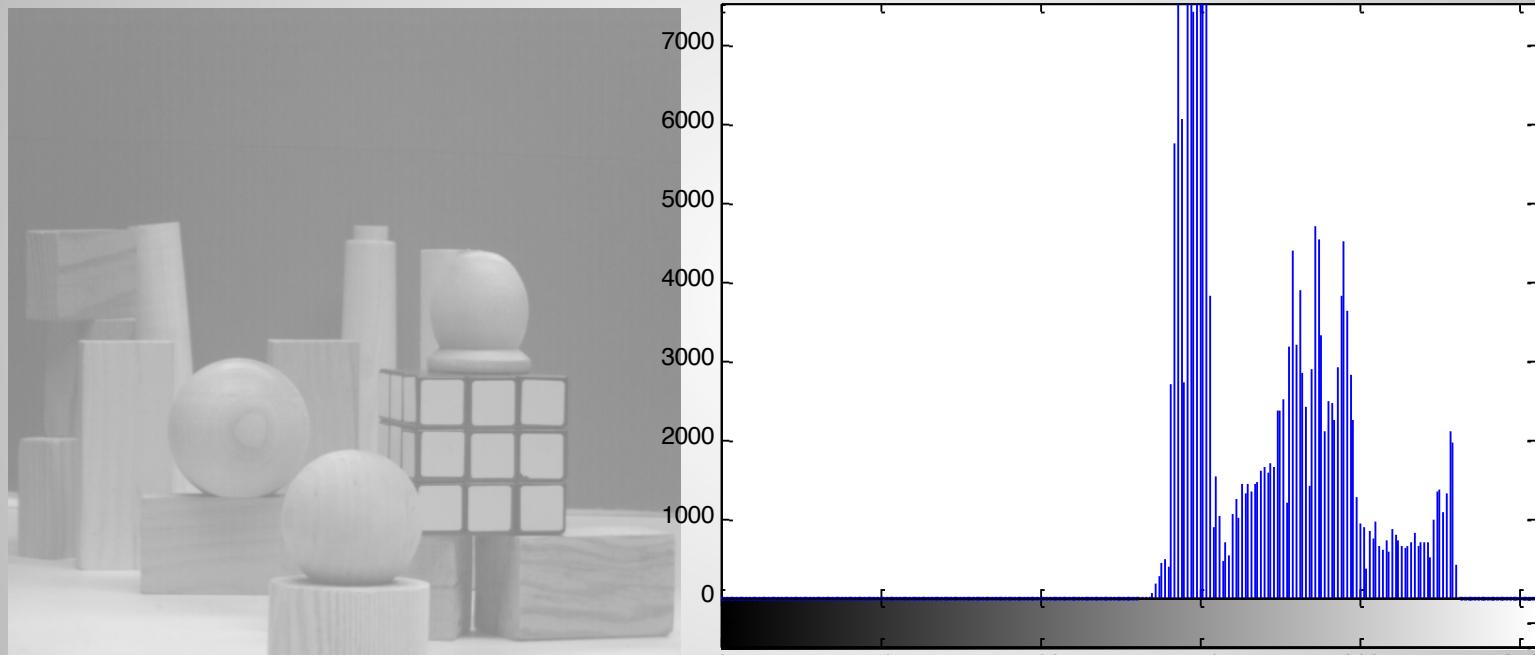


$$c=100$$

Gradation Functions

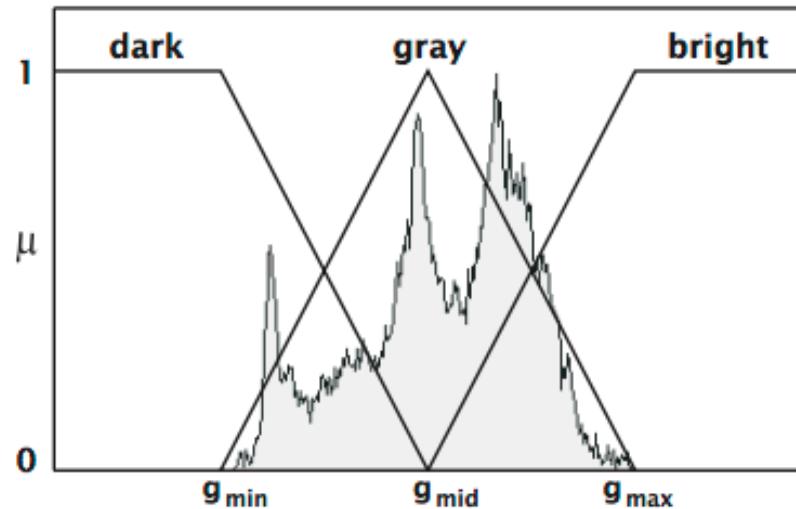
- We transform an image A into a new image A_{new} of the same size, by mapping a grey level u at pixel location p in A by a gradation function g onto a grey level $v=g(u)$.
- **Histogram Equalization:**
 - The aim is to create an image with equally distributed brightness (intensity) levels over the whole brightness (intensity) scale.
 - Enhances contrast for intensity values close to histogram maxima and decreases contrast near minima

Histogram Equalization

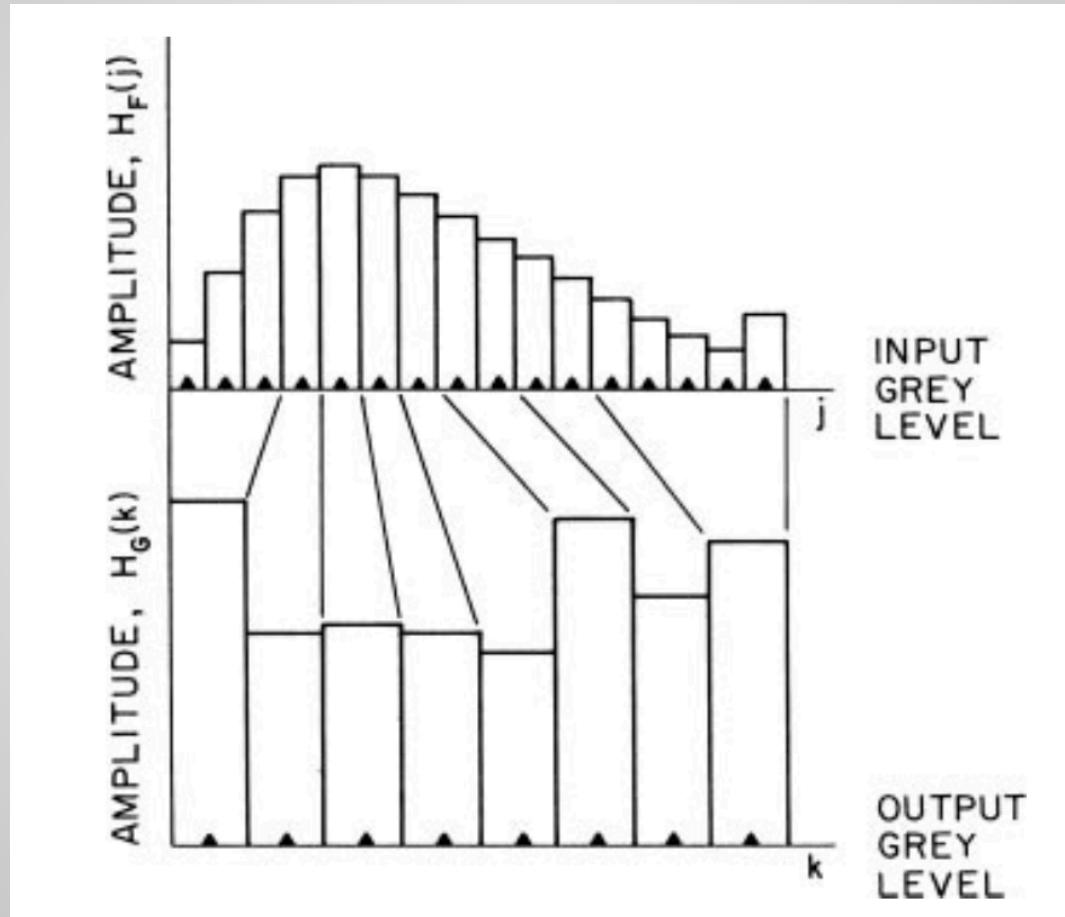


Over-exposed image

Histogram Equalization

a**b****c**

Histogram Equalization



Algorithm for Histogram Equalization

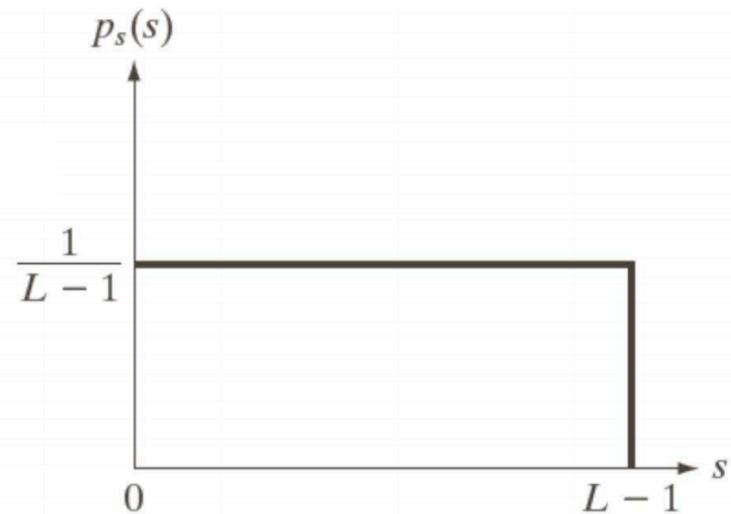
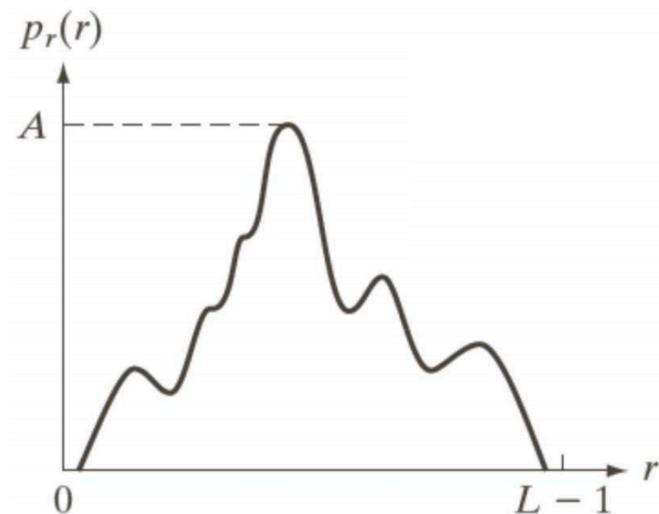
Normalized histogram p (PDF) of an image f , whose intensity values span from 0 to $L-1$

$$p_n = \frac{\text{number of pixels with intensity } n}{\text{total number of pixels}} \quad n = 0, 1, \dots, L - 1.$$

The histogram equalized image g will be defined by

$$g_{i,j} = \text{floor}\left((L - 1) \sum_{n=0}^{f_{i,j}} p_n\right),$$

Histogram Transform



Histogram Equalization – Discrete Case

- for discrete case we have:

$$\begin{aligned}
 s_k = T(r_k) &= (L-1) \sum_{j=0}^k p_r(r_j) \\
 &= \frac{(L-1)}{M \cdot N} \sum_{j=0}^k n_j \quad 0 \leq k \leq L-1
 \end{aligned}$$

- assume $L = 8$, $M = N = 64$, $M \cdot N = 4096$

r_k	n_k	$p_r(r_k) = n_k/MN$
$r_0 = 0$	790	0.19
$r_1 = 1$	1023	0.25
$r_2 = 2$	850	0.21
$r_3 = 3$	656	0.16
$r_4 = 4$	329	0.08
$r_5 = 5$	245	0.06
$r_6 = 6$	122	0.03
$r_7 = 7$	81	0.02

$$s_0 = T(r_0) = 7 \cdot \sum_{j=0}^0 p_r(r_j) = 7 \cdot p_r(0) = 1.33 \rightarrow 1$$

$$s_1 = T(r_1) = 7 \cdot \sum_{j=0}^1 p_r(r_j) = 7 \cdot (p_r(0) + p_r(1)) = 3.08 \rightarrow 3$$

$$s_2 = T(r_2) = 7 \cdot \sum_{j=0}^2 p_r(r_j) = 7 \cdot (p_r(0) + p_r(1) + p_r(2)) = 4.55 \rightarrow 5$$

$$s_3 = 5.67 \rightarrow 6, s_4 = 6.23 \rightarrow 6, s_5 = 6.65 \rightarrow 7, s_6 = 6.86 \rightarrow 7, s_7 = 7.00 \rightarrow 7$$

Histogram Equalization - Discrete

- final transform:

$$r_0 \rightarrow s_0 = 1 \Rightarrow 790 \text{ pixels map to } 1$$
$$r_1 \rightarrow s_1 = 3 \Rightarrow 1023 \text{ pixels map to } 3$$
$$r_2 \rightarrow s_2 = 5 \Rightarrow 850 \text{ pixels map to } 5$$
$$r_3 \rightarrow s_3 = 6 \Rightarrow 656 + 329 = 985 \text{ pixels map to } 6$$
$$r_4 \rightarrow s_4 = 6 \Rightarrow 656 + 329 = 985 \text{ pixels map to } 6$$
$$r_5 \rightarrow s_5 = 7 \Rightarrow 245 + 122 + 81 = 458 \text{ pixels map to } 7$$
$$r_6 \rightarrow s_6 = 7 \Rightarrow 245 + 122 + 81 = 458 \text{ pixels map to } 7$$
$$r_7 \rightarrow s_7 = 7 \Rightarrow 245 + 122 + 81 = 458 \text{ pixels map to } 7$$

Histogram Equalization Example 1

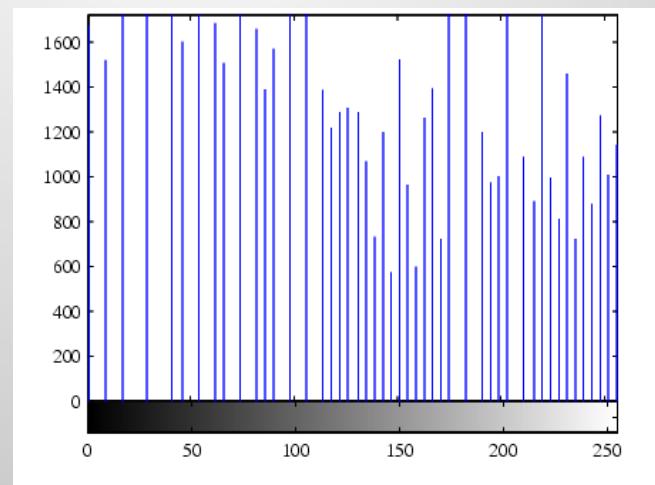
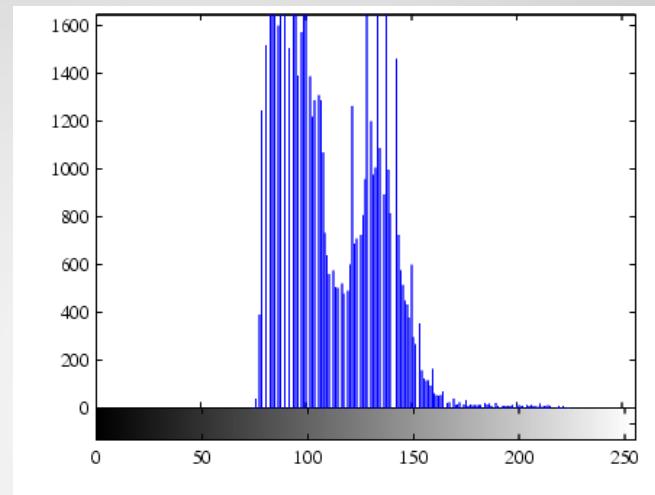


before

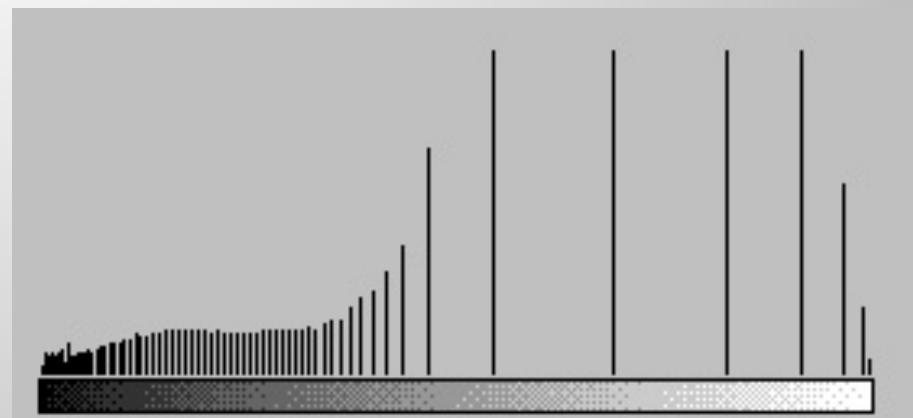
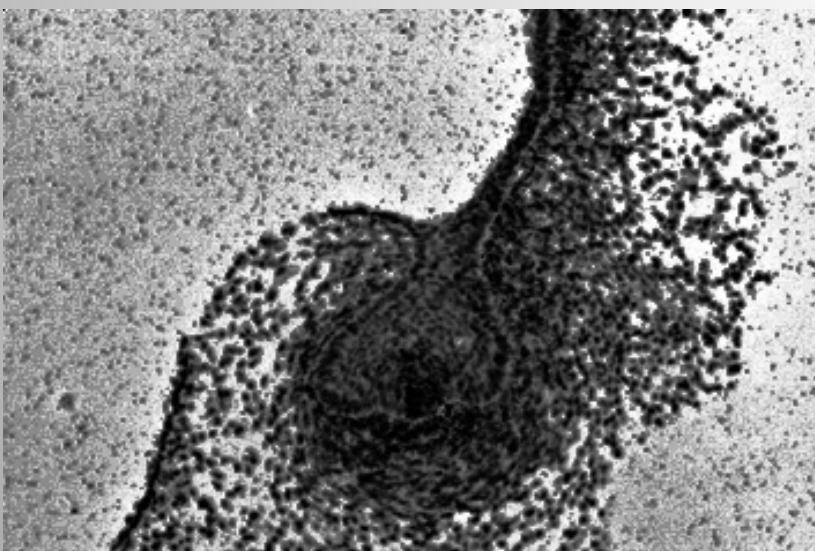
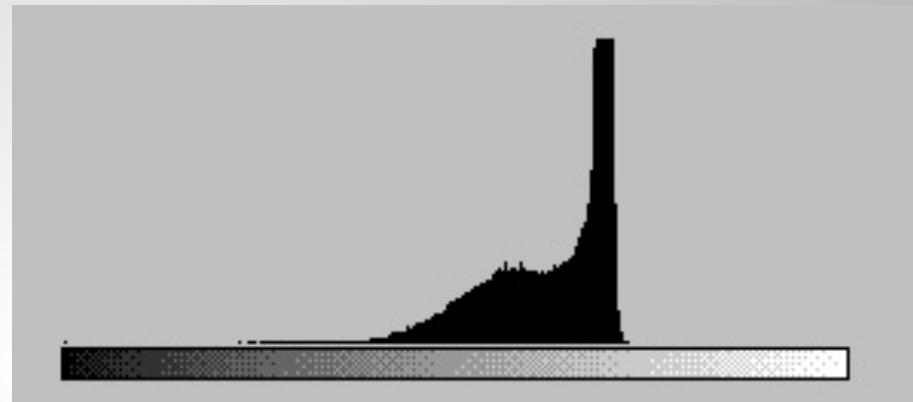
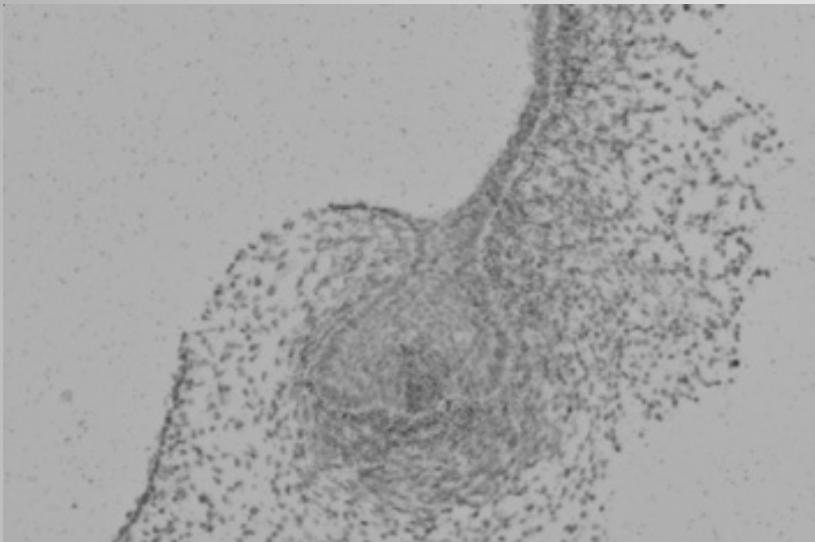


after

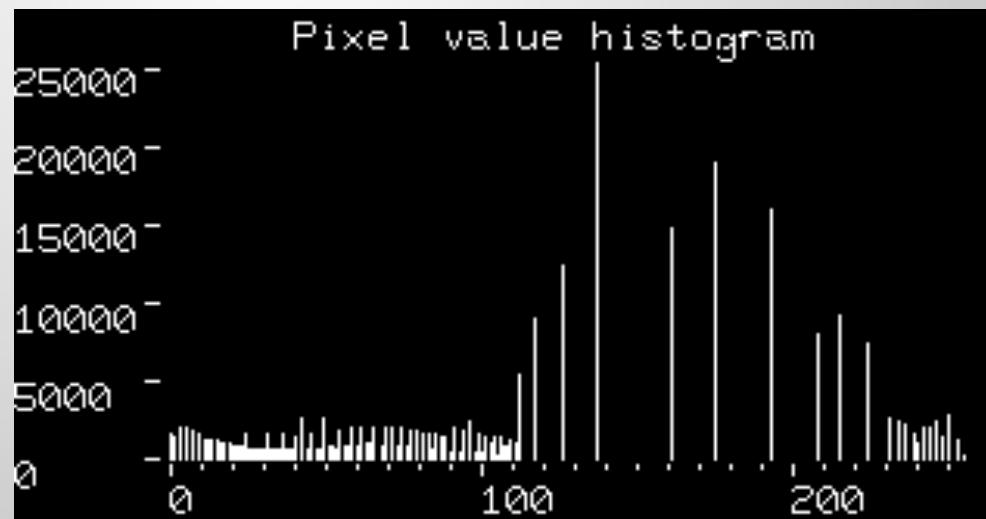
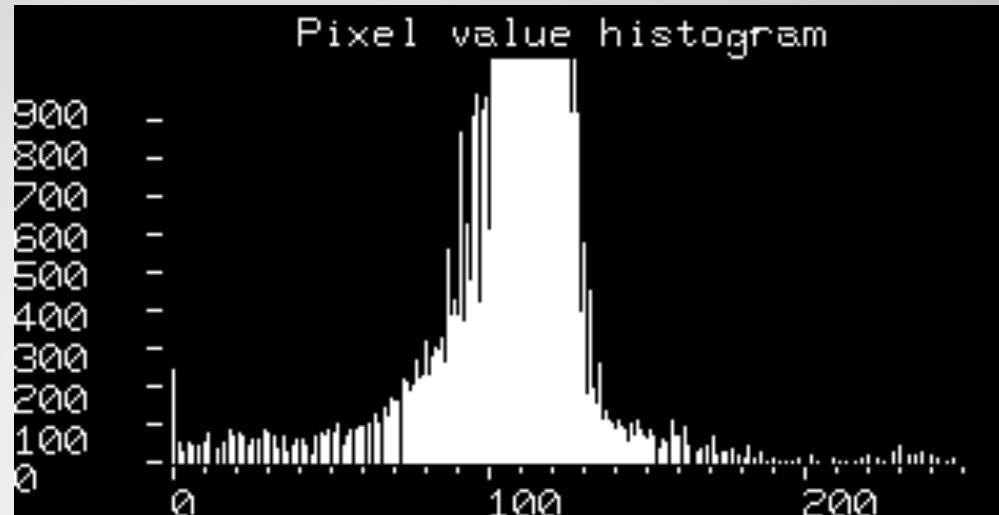
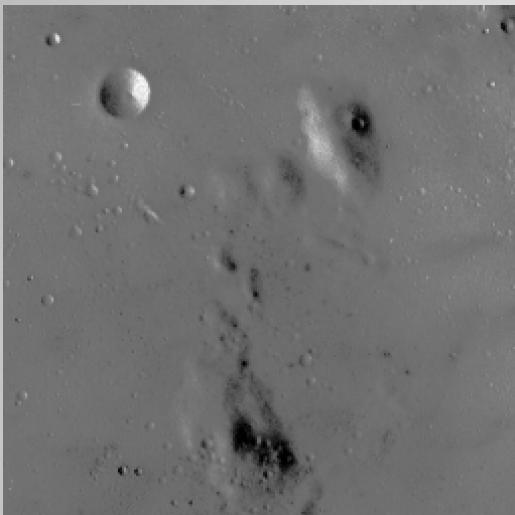
Histogram Equalization Example 2



Histogram Equalization Example 3



Histogram Equalization Example 4



Histogram Equalization Example 5



Histogram Equalization Example 6



Summary

- Edge Detection
- Noise and Smoothing
- Canny Edge Detector
- Histograms
- Histogram Equalization

Programming Assignment #1

- Image display and filtering
- **Deadline:** September 10, 2015.
- REMINDERS
 - Deadline for PA0 (Bonus) is 3rd September, 2015.
 - Submit online.